EK-12 Fairness, Bias and Audit Guidelines

Document ID: EK-12 | Category C: Evaluation knowledge base

1. Document metadata

Field	Value
Document ID	EK-12
Filename	EK-12_Fairness_Bias_and_Audit_Guidelines.pdf
Category	C - Evaluation knowledge base
Version	1.0
Creation date	2026-10-14
Data status	Factual guidance with synthetic examples
Main topics	Fairness; Bias; Audit; Privacy
Relevant skills and technologies	Ethics
Relevant entity IDs	EVA-006, QUA-004, QUA-006

Purpose and scope: define fairness rules, bias sources, audit steps and human oversight for AI proposal evaluation.

2. Executive summary

Work-related evidence only, identical formulas for all proposals, explainable scores and a client who keeps the final decision. Bias sources are paired with mitigations.

3. Principles

- Evaluate work-related evidence only.
- Apply the same formulas and thresholds to every proposal.
- Explain every score so a person can contest it.
- Keep the client in control of the final decision.

4. Characteristics that must not be used

Nationality, ethnicity, gender, age, religion, disability, family status, appearance, accent, name, photograph and location are not scored. Location or time-zone overlap may be considered only when the client states a required overlap as a requirement.

5. Sources of bias in data

Source	Example (fictional)	Mitigation
Review bias	Vague comment 'felt slow' (REV-018)	Treat vague reviews as low information (QUA-004)
Documentation bias	Freelancer with few documented projects	Report limited evidence, not low skill
Language bias	Proposal written in simple English	Do not score writing style
Popularity bias	More reviews give a more confident score	Shrink small samples toward a neutral prior
Recency bias	Old projects ignored	Use a recency factor, not exclusion

6. Auditing procedure

- Log inputs, formulas, weights and sources for each recommendation.
- Review rankings periodically for unexplained differences between groups of freelancers.
- Check that explanations never mention excluded characteristics.
- Test with synthetic cases that differ only in an irrelevant attribute; scores must be identical.

7. Handling disagreement and feedback

Clients and freelancers can challenge a score. The system shows the inputs and evidence, accepts corrected evidence and recalculates. A correction creates a new version; old versions are kept for audit.

8. Privacy and data minimization

- Store only data needed for evaluation.
- Do not send personal identifiers to external services without approval.
- Allow removal of data on request according to platform policy.

9. Human oversight

The recommendation supports the client's judgement. The interface must let the client change weights, override rankings and see flags. The assistant must not pressure a client toward the top-ranked proposal.

10. Limitations

The knowledge base is synthetic. Real deployments need real audits and legal review.

11. Summary of key findings

- Protected and irrelevant characteristics are never scored.
- Audit logs and counterfactual tests detect bias.
- Humans keep the final decision.

12. Keywords and alternative terminology

fairness; bias; audit; ethics; privacy; protected characteristics; human oversight; counterfactual test